FENIL GAJJAR

# AWS REAL TIME DAILY TASKS

# AWS S3 SERVICE

**COMPREHENSIVE GUIDE** 

THEORY + PRACTICAL

WITH REAL TIME PROJECT

**FOLLOW FOR MORE DAILY TASKS** 

fenilgajjar.devops@gmail.com

linkedin.com/in/fenil-gajjar



#### Welcome to the Ultimate Guide on

## AWS S3! 🚀

Amazon Simple Storage Service (AWS S3) is one of the most powerful and widely used cloud storage solutions. Whether you are a beginner exploring cloud storage or a professional looking to enhance your AWS skills, this guide is designed to provide you with a **comprehensive** understanding of AWS S3—from **core concepts** to **real-world implementations**.

#### What You'll Learn in This Guide

- In-Depth Theory Understand the fundamentals of AWS S3, including storage classes, security mechanisms, lifecycle policies, and more.
- ✓ Hands-On Practical Labs Step-by-step implementation of AWS S3 features to solidify your understanding.
- Real-World Project Apply your knowledge to a practical project, where you'll set up and configure S3 for a real-world use case with best practices.

By the end of this guide, you'll not only master AWS S3 theoretically but also gain hands-on experience in implementing it effectively. Get ready to dive deep into one of the most essential cloud storage services!

#### **Amazon S3 (Simple Storage Service)**

#### **Introduction to AWS S3**

Amazon Simple Storage Service (Amazon S3) is a **scalable**, **durable**, **and highly available** cloud storage service offered by AWS. It is designed to store and retrieve **any amount of data from anywhere on the web**, making it ideal for a variety of use cases, including backups, big data analytics, machine learning, media storage, and more.

AWS S3 provides **object storage**, meaning data is stored as objects inside **buckets** (containers for storing objects). Unlike traditional file systems, S3 does not use directories but instead employs a **flat structure with key-value pairs** for object organization.

#### **Key Features of AWS S3**

#### Scalability

- S3 can automatically scale storage capacity without any limitations.
- It supports **high request rates** for both read and write operations.

#### Durability & Availability

- Provides 99.99999999 (11 nines) durability, meaning your data is highly secure against loss.
- Ensures high availability with a 99.99% uptime SLA (Standard Tier).

#### **✓** Storage Classes

S3 offers different storage classes to optimize cost based on data access patterns:

- **S3 Standard** Ideal for frequently accessed data.
- S3 Intelligent-Tiering Automatically moves data between storage classes to optimize costs.
- S3 Standard-IA (Infrequent Access) For data that is accessed less frequently but needs rapid retrieval.
- S3 One Zone-IA A lower-cost option for infrequent access data stored in a single AWS Availability Zone.
- S3 Glacier Low-cost storage for archival purposes with retrieval times from minutes to hours.
- S3 Glacier Deep Archive The cheapest storage for long-term data retention (retrieval takes up to 12 hours).

#### ✓ Security & Compliance

- Data Encryption Supports server-side encryption (SSE) and client-side encryption.
- IAM Policies & Bucket Policies Fine-grained access control using Identity and Access Management (IAM).
- Versioning Maintains multiple versions of an object to prevent accidental deletion.
- MFA Delete Adds an extra layer of security by requiring Multi-Factor
   Authentication (MFA) for deletions.

#### Data Lifecycle Management

Automates data movement across storage classes using Lifecycle Policies.

 Allows you to expire or delete old objects automatically, reducing storage costs.

#### Performance & Speed

- Parallel Uploads & Multipart Uploads for handling large files efficiently.
- Transfer Acceleration Speeds up data uploads using AWS global edge locations.

#### **V** Event Notifications & Integrations

- Triggers events (e.g., file uploads, deletions) using AWS Lambda, SNS, or SQS for automation.
- Seamlessly integrates with AWS CloudFront, AWS Athena, AWS Glue, and more.



#### Life Before AWS S3: The Struggles of

#### **Traditional Storage**

Before AWS S3 revolutionized cloud storage, businesses and individuals relied on traditional storage methods like on-premises data centers, local servers, external hard drives, and network-attached storage (NAS). While these solutions worked for decades, they had significant limitations and challenges that made data management complex, costly, and unreliable.

#### 🚧 The Struggles of Traditional Storage Systems

#### **Limited Scalability**

- Storage expansion required purchasing and installing new physical hardware.
- Companies had to predict future storage needs, often leading to underutilization or overinvestment.
- Once storage was full, there was no on-demand expansion—manual intervention was always needed.

#### High Costs & Maintenance Overhead

- Businesses had to invest in expensive servers, cooling systems, power backup, and dedicated IT staff.
- Data center maintenance included hardware repairs, upgrades, and space management.

 High initial investment and ongoing operational costs made storage management financially draining.

#### Risk of Data Loss & Limited Durability

- Single-point failures like hard disk crashes, power failures, or accidental deletions could result in permanent data loss.
- Organizations had to create manual backup strategies, often requiring
   offsite backup storage to prevent catastrophic failures.
- Data replication was complicated and required additional storage infrastructure.

#### **X** Complex Backup & Disaster Recovery

- Companies relied on tape backups, RAID setups, and offsite data centers for disaster recovery.
- Backups needed to be scheduled manually, making recovery slow and unreliable.
- If a data center was damaged due to a natural disaster, fire, or cyberattack,
   data restoration was time-consuming and expensive.

#### Security Challenges & Data Breaches

- Physical servers were vulnerable to hardware theft, unauthorized access,
   and internal threats.
- Organizations had to implement complex security policies, including firewalls, VPNs, and manual encryption.
- Managing user permissions and access control was a manual process prone to misconfigurations.

#### Limited Accessibility & Remote Collaboration

- Data stored on local servers was only accessible within office premises.
- Employees working remotely or in different locations struggled with slow
   VPN connections and file transfer limitations.
- Large files had to be shared via external hard drives, FTP servers, or email attachments, slowing down workflows.

#### Manual Storage Management & Lack of Automation

- IT teams had to monitor storage usage, delete old files manually, and migrate data to different storage systems.
- Archiving old data was time-consuming and inefficient without automated lifecycle policies.
- Data retrieval times depended on network bandwidth and server capacity,
   often leading to slow performance.

# Life After AWS S3: The Future of Scalable, Secure & Cost-Effective Storage

With the introduction of **AWS S3**, cloud storage underwent a revolutionary transformation. Businesses, developers, and enterprises no longer need to worry about **hardware limitations**, **data loss**, **or high maintenance costs**. Instead, they can focus on **innovation**, **agility**, **and seamless data management** while AWS takes care of storage scalability, security, and durability.

#### The Advantages of AWS S3 Over Traditional Storage

#### Unlimited Scalability

- No need to purchase, install, or maintain physical storage.
- Store **unlimited data** without worrying about running out of space.
- Easily scale up or down based on real-time storage needs without manual intervention.

#### 💸 Cost-Effective & Pay-as-You-Go Pricing

- No upfront capital investment—you pay only for the storage you use.
- Multiple storage classes allow businesses to optimize costs by moving infrequently accessed data to cheaper tiers (e.g., S3 Glacier for long-term archival).
- No need to maintain costly data centers, cooling, or power infrastructure.

#### High Durability & Data Protection (11 Nines Reliability)

- AWS S3 ensures 99.999999999 durability by replicating data across multiple availability zones.
- Features like versioning, cross-region replication, and automatic backups prevent accidental data loss.
- Object lock allows for immutable data storage, preventing unauthorized modifications or deletions.

#### 🌎 Global Accessibility & Seamless Collaboration

- Data stored in S3 is accessible from anywhere in the world via APIs, AWS SDKs, and CLI.
- Remote teams can collaborate effortlessly without the need for VPNs or complex file-sharing systems.
- Integrates with AWS services like Lambda, EC2, CloudFront, and RDS,
   making data storage part of a larger cloud ecosystem.

#### **K** Fully Managed & Automated Operations

- No manual maintenance—AWS handles hardware failures, backups, and security patches.
- Lifecycle policies allow for automatic archiving, deletion, and storage class transitions.
- Event-driven architecture: AWS S3 triggers AWS Lambda functions,
   enabling serverless automation for data processing.

#### **Advanced Security & Compliance**

- Built-in encryption at rest and in transit ensures maximum data protection.
- IAM policies, bucket policies, and access control lists (ACLs) enable
   fine-grained access control.
- Supports compliance standards like GDPR, HIPAA, and SOC 2, making it ideal for enterprise-grade security.

#### High Performance & Fast Data Retrieval

- Data retrieval is **instantaneous** compared to traditional backup solutions.
- S3 Select and Glacier Select allow for querying data directly within S3,
   reducing data transfer and compute costs.
- Integrated with AWS CloudFront (CDN) for faster global content delivery.

### **% When Should You Use AWS S3?**

AWS S3 is a **versatile**, **scalable**, **and cost-effective** storage solution, but when exactly should you use it? Whether you are an **individual developer**, **a startup**, **or an enterprise**, S3 provides **reliable storage** for various use cases. Let's explore in detail when and why you should use AWS S3.

#### A

#### **Storing & Serving Static Website Content**

#### When to Use?

- When hosting static websites, blogs, or landing pages.
- When serving HTML, CSS, JavaScript, images, and videos directly from the cloud.
- When you need a cost-effective alternative to traditional web hosting.

#### Why Use S3?

- Supports static website hosting with a public endpoint.
- Works seamlessly with AWS CloudFront (CDN) for faster global content delivery.
- Eliminates the need for a **dedicated web server**, reducing costs.

#### **Example**:

A company hosting a **product landing page** with HTML, CSS, and JS files can use **AWS S3** instead of **deploying a web server**, making the setup **simpler**, **cheaper**, and **highly available**.

#### Data Backup & Disaster Recovery

#### When to Use?

- When you need to store backups of databases, files, and system snapshots.
- When implementing a disaster recovery strategy for critical business data.
- When **long-term archival** of data is required.

#### Why Use S3?

- 11 nines durability (99.99999999%) ensures data is never lost.
- Versioning & cross-region replication protect against accidental deletions or regional failures.
- Lifecycle policies automatically move older backups to cheaper storage classes like S3 Glacier.

#### **Example**:

A financial firm needs to store client transaction records for 7+ years. By using S3 Standard for recent data and S3 Glacier for long-term retention, they achieve cost savings and compliance.

#### Big Data Storage & Analytics

#### When to Use?

• When storing massive datasets for Al, ML, and analytics.

- When needing fast access to unstructured data (logs, IoT data, sensor data).
- When working with AWS services like Athena, Redshift, and Glue for analytics.

#### Why Use S3?

- S3 Select & Glacier Select allow you to query data directly within S3,
   reducing processing costs.
- Integrates with AWS Data Lakes, AI/ML, and analytics tools.
- Supports parallel processing for faster computations.

#### **Example**:

An e-commerce company **analyzes customer behavior** by storing **clickstream data in S3** and running queries using **AWS Athena**, enabling real-time insights.

#### Media Hosting & Content Distribution

#### When to Use?

- When storing and streaming videos, images, and large media files.
- When distributing software updates, gaming assets, or digital content globally.
- When requiring **fast and reliable delivery** of large files.

#### Why Use S3?

- CloudFront integration ensures low-latency content delivery worldwide.
- Supports **multipart uploads** for handling large files efficiently.
- Pay-as-you-go pricing makes it cost-effective for media storage.

#### **Example**:

A video streaming platform uses **AWS S3 to store high-resolution videos** and serves them via **CloudFront for faster playback** across different devices.

#### Hosting Machine Learning (ML) & Al Training Data

#### When to Use?

- When storing datasets for Al/ML model training.
- When needing high availability and fast retrieval of training data.
- When working with TensorFlow, PyTorch, or SageMaker.

#### Why Use S3?

- Supports high-throughput data access for ML training models.
- Integration with AWS SageMaker for seamless training workflows.
- Versioning support helps track changes in training datasets.

#### **Example**:

A self-driving car company stores **terabytes of video footage** in S3 and uses **AWS SageMaker** to train machine learning models on this data.

#### Storing Logs & Application Data

#### When to Use?

- When storing server logs, API request logs, and application metrics.
- When integrating with AWS Lambda for event-driven workflows.

• When needing a **centralized logging solution** for monitoring.

#### Why Use S3?

- Supports real-time log storage from EC2, Lambda, and CloudTrail.
- Lifecycle policies help archive old logs to S3 Glacier.
- Works with AWS Athena for querying logs without loading them into a database.

#### **Example**:

A DevOps team stores **AWS CloudTrail logs in S3** and queries them using **Athena** to analyze security incidents.

#### Software Development & DevOps Pipelines

#### When to Use?

- When storing Docker images, deployment artifacts, and build logs.
- When integrating with CI/CD pipelines for automated deployments.
- When needing a **central storage for configuration files**.

#### Why Use S3?

- Integrates with AWS CodePipeline, CodeBuild, and Jenkins.
- Versioning support ensures that old deployment artifacts are not lost.
- IAM policies allow developers to control access securely.

#### **Example**:

A DevOps team stores **Terraform state files in S3** to maintain infrastructure consistency across deployments.

#### IoT Data Storage & Processing

#### When to Use?

- When collecting sensor data from IoT devices.
- When performing real-time analytics on IoT data.
- When needing a scalable and durable storage solution.

#### Why Use S3?

- Works with AWS IoT Core and AWS Lambda for real-time event processing.
- Automates lifecycle policies to move old data to cost-effective storage.
- Handles petabytes of sensor data efficiently.

#### **Example**:

A smart home company stores **temperature and humidity data** from IoT sensors in **AWS S3** and triggers alerts using **AWS Lambda**.

### Nhen Should You Not Use AWS S3?

While **AWS S3** is a powerful and versatile storage solution, it is *not* the best fit for every use case. There are scenarios where using **alternative AWS services** or a **different storage approach** would be more efficient. Here's when you should avoid using AWS S3 and consider other options instead.

#### High-Performance, Low-Latency Databases Needed

AWS S3 is an **object storage service**, meaning it is optimized for **large-scale storage and retrieval of files**, not for **frequent**, **low-latency transactions**. If your application requires **millisecond response times** or **real-time database queries**, you should consider:

- Amazon RDS (Relational Database Service) for structured, transactional data.
- Amazon DynamoDB for fast, NoSQL key-value storage.
- Amazon ElastiCache (Redis/Memcached) for ultra-low-latency data caching.

**Example**: An e-commerce website that needs **instant product catalog lookups** should use **DynamoDB or RDS**, not S3.

#### **H** Running an Operating System or Hosting a Database

AWS S3 is **not a block storage system** and cannot be used to run **operating systems or databases** directly. If you need **persistent storage** for a server or database, use:

Amazon EBS (Elastic Block Store) for high-performance, persistent storage

attached to EC2.

- Amazon EFS (Elastic File System) for scalable, shared file storage.
- **Example**: A company running a **MySQL** database should use **Amazon RDS** with EBS, not S3.

#### Hosting Dynamic Websites or Applications

S3 can serve **static websites** but cannot process **dynamic content, handle authentication, or run server-side logic**. If you need a **full-fledged web application**, use:

- AWS Elastic Beanstalk or EC2 for hosting dynamic web applications.
- AWS Lambda & API Gateway for serverless applications.
- Amazon Lightsail for easy, managed web hosting.
- Example: A social media platform that requires user logins, real-time interactions, and database queries should use EC2 or AWS Lambda, not S3.

# Need File System Features Like Hierarchical Structure & File Locks

AWS S3 is **object storage**, meaning it does not support:

- Traditional folder-based hierarchical structures (though prefixes can be used).
- File locking, preventing users from editing files concurrently.
- POSIX-compliant file operations needed for applications requiring direct file system access.

If you need a shared, mountable file system, consider:

- Amazon EFS for scalable, shared file access across multiple instances.
- Amazon FSx (for Windows or Lustre) for enterprise-grade file storage.

**Example**: A company managing video editing projects that require team collaboration on large files should use Amazon FSx, not S3.

#### **Solution** Cost-Sensitive Frequent Small File Access

S3 charges for **API requests**, which means frequent **small file access** (e.g., retrieving thousands of tiny logs per second) can lead to unexpected costs. If your workload requires frequent read/write operations, consider:

- Amazon EFS (if shared file storage is needed).
- Amazon DynamoDB (for structured, high-speed access).
- Amazon S3 Intelligent-Tiering (to optimize storage costs automatically).
- **Example**: A real-time log processing system fetching millions of small log files should use EFS or DynamoDB, not S3.

#### Complex Transactions & Record Updates

AWS S3 does not support **transactions**, **record locking**, **or atomic updates** like a relational database. If you need:

- ACID transactions for banking, financial systems, or inventory tracking → use
   Amazon RDS (MySQL, PostgreSQL, etc.).
- Frequent updates to structured records → use Amazon DynamoDB or Aurora.

**Example:** A **banking system** that updates account balances after each transaction should use Amazon RDS or DynamoDB, not S3.

#### Storing Sensitive Data Without Proper Security Controls

By default, S3 buckets are private, but misconfigurations can lead to data breaches. If your use case involves highly sensitive or regulated data, ensure:

- IAM Policies, Encryption, and VPC Endpoints are properly configured.
- Use AWS KMS (Key Management Service) for encryption.
- Consider Amazon Macie to detect security risks in your S3 buckets.
- **Example:** A healthcare company storing patient medical records must encrypt data at rest and in transit and apply strict IAM permissions to prevent exposure.



### **X** How AWS S3 Works? – Deep Dive into Its

#### **Architecture**

AWS Simple Storage Service (S3) is a scalable, durable, and highly available **object storage service** that allows users to store and retrieve data anytime, anywhere. Understanding its architecture is crucial to knowing how S3 handles storage, security, and performance. Let's break it down in detail.



#### T S3 Architecture Overview

AWS S3 follows a distributed architecture designed for high availability, durability, and scalability. It consists of multiple components that work together to ensure seamless data storage and retrieval.



#### S3 Buckets

The **S3 bucket** is a **logical container** for storing objects (files).

- Every bucket has a globally unique name within AWS.
- Data stored in a bucket is replicated across multiple AWS Availability Zones (AZs) for durability.
- Buckets can be private or public depending on permissions.
- Each bucket resides in a specific AWS region chosen by the user.
- **Example**: A company storing user-uploaded profile pictures would create an **S3** bucket named my-app-profile-images.

#### **Welliam Strain Strain Strain Strain Welliam Strain Stra**

S3 stores **objects**, which are files along with their metadata.

- Objects are stored inside buckets.
- Each object has a unique key (filename) + metadata.
- Objects can be any size (from a few KBs to 5TB).
- S3 provides **Versioning**, allowing users to keep multiple versions of the same object.
- **Example**: Storing profile-pic.png in my-app-profile-images bucket results in an object like:
- s3://my-app-profile-images/profile-pic.png

#### Object Identifiers (Key + Version ID)

Each object in S3 is uniquely identified by:

- Bucket Name (e.g., my-app-profile-images).
- Object Key (File Path/Name) (e.g., profile-pic.png).
- Version ID (if versioning is enabled).
- **Example**:
- s3://my-app-profile-images/user123/profile-pic-v1.png

#### S3 Storage Classes (Data Tiers)

S3 offers **multiple storage classes** for different use cases:

- Standard → For frequent access, high durability.
- Intelligent-Tiering → Auto-moves objects between access tiers based on

usage.

- Standard-IA (Infrequent Access) → Lower-cost storage for less frequently accessed data.
- One Zone-IA → Similar to IA but stored in a single AZ (less redundancy).
- Glacier & Glacier Deep Archive → For long-term data archiving with retrieval delays.
- **Example**: A backup system might store **daily logs** in Standard and **old logs** in Glacier.

#### T Deep Dive into S3 Working Architecture

#### 1 Data Ingestion & Storage

When a user uploads a file to S3:

- The file is divided into multiple parts (for large files).
- Each part is **encrypted (if enabled)** before leaving the user's device.
- S3 distributes the data across multiple storage nodes in different Availability
   Zones (AZs).
- S3 returns an ETag (Entity Tag) for verification, ensuring data integrity.
- **Example**: Uploading a **1GB video** results in S3 splitting it into smaller parts and storing them across multiple AZs.

#### **2** Retrieval & Access Mechanism

When a user retrieves a file from S3:

S3 fetches the file parts from distributed nodes.

- It reassembles them into a complete file.
- The file is returned via HTTP(S) using a unique URL.
- S3 optimizes retrieval with range requests, allowing partial file downloads.
- **Example**: A media streaming service can use **range requests** to **fetch only a portion** of a video file for buffering.

#### **3** Security & Access Control

AWS S3 provides multiple security layers:

- **IAM Policies (Identity & Access Management)**
- Control access at user level (who can read/write objects).
- Bucket Policies & ACLs (Access Control Lists)
- Define public or private access for entire buckets or individual files.
- Encryption
- Server-side (SSE-S3, SSE-KMS, SSE-C) → Data is encrypted automatically.
- Client-side → User encrypts data before uploading.
- **Example**: A company storing **customer invoices** can enable **SSE-KMS** encryption to secure sensitive data.

#### **14** Lifecycle Policies & Data Management

AWS S3 allows automated **object lifecycle management**:

- Move objects to a lower-cost storage class after X days.
- Automatically delete objects after a specified period.

**Example**: Log files stored in **S3 Standard** can be moved to **Glacier** after **30** days and deleted after **1 year**.

#### **\*\*5** Performance Optimization

AWS S3 is optimized for high-speed access:

- Multipart Uploads → Increases upload speed for large files.
- Byte-Range Fetches → Allows downloading parts of files.
- S3 Transfer Acceleration → Uses AWS Edge Locations to speed up uploads.
- Replication (Cross-Region & Same-Region) → Auto-replicates objects for redundancy.
- **Example**: A global application can enable **Cross-Region Replication (CRR)** to store backups in **multiple AWS regions**.

#### How AWS S3 Handles Scalability & Reliability?

#### Scalability

- S3 automatically scales as data grows.
- There is **no limit** on the number of objects stored.
- It supports high concurrency, handling millions of requests per second.
- **Example**: Netflix stores **petabytes of video content** on **S3**, handling millions of user requests seamlessly.

#### Durability & Availability

- S3 provides **99.99999999%** (**11 9's) durability**, meaning your data is **extremely safe**.
- Data is replicated across multiple AZs, ensuring high availability.
- **Example**: Even if one AZ **fails**, S3 ensures data remains accessible from other AZs.

#### How S3 Integrates with Other AWS Services?

AWS S3 is a **core service** that integrates with multiple AWS products:

- Lambda → Run serverless functions on S3 object events (e.g., trigger a function when a file is uploaded).
- CloudFront → Deliver S3 content globally with CDN caching.
- AWS Backup → Automate backups of S3 data.
- Athena → Query S3 data using SQL without a database.
- Macie → Detect sensitive data in S3 for security compliance.
- **Example**: A website serving images worldwide can use **S3 + CloudFront** for **faster content delivery**.



#### Key Components of AWS S3 Service

AWS S3 (Simple Storage Service) is a highly scalable, durable, and secure object storage service. It consists of several key components that work together to provide seamless data storage and retrieval. Let's dive into each of them in detail.



#### S3 Buckets – The Storage Containers

A **bucket** is like a folder that stores objects (files, images, videos, backups, etc.).

- **Globally unique name:** Each bucket name must be unique across AWS.
- **Region-specific:** You choose a **region** where the bucket is stored.
- Access Control: Permissions define who can read/write data.

#### **Example:**

A company creates a bucket "mycompany-logs" in the us-east-1 region to store application logs.



#### **Objects – The Actual Data Files**

Objects are the actual **files stored inside a bucket**.

- Each object consists of data + metadata + unique key (filename).
- Objects can be up to 5TB in size.
- Stored in a flat structure (no traditional folder hierarchy, but prefixes can be used).

#### **Example:**

A profile picture (user123.jpg) is stored inside the bucket:

→ Object Key: images/profile/user123.jpg

#### Object Key – The Unique Identifier

Each object in a bucket is identified using a unique key (similar to a filename).

- Helps in retrieving data efficiently.
- Supports **prefixes for organization** (e.g., images/, logs/).

#### **Example:**

A log file stored with a structured naming convention:

- logs/2024/02/app\_log\_01.txt
- logs/2024/02/app\_log\_02.txt

#### Metadata – Descriptive Information About Objects

Metadata provides extra information about an object, such as:

- Content-Type (e.g., image/png, text/csv).
- Last modified date.
- Encryption status.
- Custom metadata (e.g., X-Tag: confidential).

#### **Example:**



#### Recursion - Security & Permissions

AWS S3 provides multiple ways to control who can access buckets & objects.

#### Bucket Policies

- JSON-based rules that define who can access what in a bucket.
- Applied at the bucket level.

#### IAM Policies

- Define permissions for **AWS users, roles, and groups**.
- Allows fine-grained access control.

#### ACLs (Access Control Lists) – Legacy Method

Used to grant specific users or accounts access to buckets or objects.

#### Pre-Signed URLs

• Temporary links to **grant time-limited access** to objects.

#### **Example:**

A website generates a pre-signed URL to allow users to download a report without making the file public.

#### S3 Storage Classes – Cost & Performance Optimization

S3 offers multiple storage classes to optimize cost and performance based on access patterns.

**S3 Standard** – Frequently accessed data.

- **S3 Intelligent-Tiering** Auto-optimizes based on usage.
- S3 Standard-IA (Infrequent Access) For rarely accessed but instantly available data.
- **S3 One Zone-IA** Cheaper, single AZ storage.
- S3 Glacier & Glacier Deep Archive Long-term, archival storage.

#### **Example:**

A media company stores **frequently accessed videos** in **S3 Standard** and moves **older videos** to **S3 Glacier** to save costs.

#### S3 Versioning – Track Changes to Objects

- Allows multiple versions of the same object to be stored.
- Helps in recovering from accidental deletions or overwrites.
- Latest version is always served, but previous versions are retained.

#### Example:

A document editing app enables versioning to keep track of old file versions.

#### S3 Lifecycle Policies – Automated Data Movement

- Automatically transitions objects between storage classes based on rules.
- Helps in **cost optimization** by moving old files to cheaper storage.

#### **Example:**

A company sets a **lifecycle rule** to:

✓ Keep files in S3 Standard for 30 days.

- ✓ Move them to S3 Standard-IA after 30 days.
- ✓ Archive them in S3 Glacier after 1 year.

#### **S3** Replication – Data Redundancy Across Regions

- Copies objects from one bucket to another across AWS Regions.
- Ensures disaster recovery & global availability.

#### **Example:**

A company **replicates** backups from us-east-1 to ap-south-1 for **disaster recovery**.

#### S3 Event Notifications – Trigger AWS Services

- S3 can send notifications when objects are created, modified, or deleted.
- Can trigger **AWS Lambda, SNS, SQS**, etc.

#### **Example:**

A **Lambda function** is triggered whenever a **new image is uploaded**, automatically processing it for a thumbnail.

#### S3 Encryption – Secure Your Data

AWS S3 supports encryption at rest and in transit to ensure data security.

**♀** Server-Side Encryption (SSE)

- SSE-S3: AWS manages encryption keys automatically.
- SSE-KMS: Uses AWS Key Management Service (KMS) for better key control.
- SSE-C: Bring Your Own Keys (BYOK) for custom security.

#### Client-Side Encryption

- Encrypts data **before uploading** to S3.
- Uses AWS KMS or user-managed encryption keys.

#### **Example:**

A healthcare company encrypts sensitive patient records using **SSE-KMS** to meet **compliance regulations (HIPAA, GDPR, etc.)**.

# S3 Cross-Region Replication (CRR) & Same-Region Replication (SRR)

AWS S3 supports data replication to improve availability and disaster recovery.

#### 

- Automatically **copies objects** from one **region** to another.
- Used for disaster recovery & global access.

#### Same-Region Replication (SRR)

- Copies objects within the same AWS region.
- Useful for compliance & logging.

#### **Example:**

A financial company replicates data from us-east-1 to eu-west-1 to meet **regulatory requirements** and ensure **geo-redundancy**.

#### S3 Object Lock – Prevent Accidental Deletions

AWS S3 **Object Lock** protects objects from being **deleted or modified**.

- **Governance Mode:** Users need special permissions to delete objects.
- Compliance Mode: Objects cannot be deleted or changed until the retention period expires.

#### **Example:**

A legal firm locks archived documents for 7 years to comply with regulatory requirements.

#### S3 Storage Lens – Storage Analytics & Insights

S3 Storage Lens provides insights & recommendations on storage usage.

- Helps identify unused objects to save costs.
- Monitors bucket activity to detect anomalies.

#### **Example:**

An enterprise uses **S3 Storage Lens** to analyze storage trends and **automate lifecycle policies** to save money.



#### S3 Multi-Part Upload – Handling Large Files

For large files (>100MB), S3 supports multi-part upload to:

- Upload parts in parallel, speeding up transfers.
- Resume uploads if interrupted.

#### **Example:**

A video streaming platform uploads large 4K videos using multi-part upload, ensuring faster and reliable uploads.

#### **S3** Transfer Acceleration – Faster Global Uploads

S3 Transfer Acceleration routes uploads through AWS Edge locations, reducing latency for global users.

#### **Example:**

A media company with users worldwide **uploads large datasets faster** from different continents.

#### S3 Requester Pays – Shifting Download Costs

- By default, **bucket owners pay** for requests.
- With Requester Pays, users who download objects pay the transfer costs.

# **Example:**

A government agency hosts public datasets, and users who download the data bear the costs instead of the agency.

# S3 Access Points – Simplified Access Management

- Create different access points for various applications.
- Helps in controlling permissions per department.

# **Example:**

A company with multiple teams creates separate access points for finance, HR, and IT teams using the same bucket.

# 🎭 S3 Object Tagging – Organizing & Managing Objects

- Assign custom metadata tags to objects.
- Helps in cost allocation, search, and policy enforcement.

# Example:

An e-commerce company tags objects as "Product-Images", "Invoices", and "Customer-Data" for easier tracking.

# S3 Batch Operations – Automating Large-Scale Tasks

Perform operations on millions of objects with a single API request.

 Modify object properties, copy objects, delete files, or restore from Glacier.

# **Example:**

A company moves 100,000 files from one storage class to another using S3 Batch Operations.

# **S3** Object Lambda – Dynamic Data Processing

- Modifies data on the fly before returning it to the user.
- Eliminates the need to store multiple versions of an object.
- Uses AWS Lambda to process objects dynamically.

# **Example:**

A media company dynamically watermarks images before delivering them to users, without modifying the original image in S3.

# S3 Select & Glacier Select – Query Data Inside Objects

- Allows you to run SQL queries on CSV, JSON, and Parquet files stored in S3.
- Reduces data transfer costs by retrieving only the needed data.
- S3 Select works on S3 standard, while Glacier Select works on archived data.

# **Example:**

A data analytics team runs **SQL** queries on large CSV logs stored in S3, retrieving only **specific rows** instead of downloading the entire file.

# S3 Intelligent-Tiering – Automatic Cost Optimization

- Moves objects between different storage classes automatically based on access patterns.
- No retrieval fees or performance impact.
- Ideal for data with unpredictable access patterns.

# **Example:**

A marketing team stores customer analytics data, which S3 automatically moves to lower-cost tiers when it's not accessed frequently.

# S3 Block Public Access – Prevent Accidental Exposure

- Prevents buckets and objects from being publicly accessible.
- Applies at the bucket or account level.
- Ensures data security and compliance.

# **Example:**

A financial company ensures that sensitive documents cannot be publicly accessed, even if someone mistakenly applies public permissions.



#### **S3 Event Notifications – Automate Workflows**

- Triggers Lambda functions, SQS, or SNS when new objects are added, modified, or deleted.
- Helps automate real-time processing.

# **Example:**

A video-sharing platform triggers an AWS Lambda function to compress and optimize videos whenever a new video is uploaded to S3.

# S3 Inventory – Track and Audit Objects

- Provides a detailed list of all objects in a bucket.
- Helps with security audits, cost management, and compliance tracking.
- Can generate reports daily or weekly.

# **Example:**

A compliance officer runs an S3 Inventory Report to ensure that all critical objects have proper encryption and lifecycle policies applied.



# When Should We Use Each AWS S3

# **Storage Class?**

AWS S3 offers **multiple storage classes**, each optimized for **different use cases** based on access frequency, cost, and durability requirements. Let's explore when to use each storage class in detail.

# S3 Standard – High Availability & Performance

Best For: Frequently accessed data

Durability: 99.999999999% (11 9's)

Availability: 99.99%

Cost: Highest among all classes

# **Use Cases:**

- ✓ Websites storing dynamic content (images, videos, documents).
- ✓ Frequently accessed logs, reports, and datasets.
- ✓ Mobile apps storing user-generated content (profile pictures, posts).
- ✓ Al/ML training datasets that require fast and frequent access.
- **Example:** An e-commerce website serving **product images & videos** in real time.

# **11** 2 S3 Intelligent-Tiering – Cost Optimization for

# **Unpredictable Access**

Best For: Data with changing access patterns

Durability: 99.999999999 (11 9's)

Availability: 99.9%

Cost: Slightly lower than S3 Standard (auto-optimizes based on usage)

#### **Use Cases:**

- ✓ Data that might be accessed frequently **now**, but infrequently **later**.
- ✓ **Data lakes** where access frequency is uncertain.
- ✓ Al/ML datasets where some files are accessed frequently, and others rarely.
- ✓ SaaS applications storing customer data with changing usage patterns.

**Example:** A social media platform storing user photos & videos, where **some** are accessed daily, and others rarely.

# **3**S3 Standard-IA (Infrequent Access) – Lower Cost for Less Used Data

- Best For: Data that is accessed less frequently but needs fast retrieval
- Durability: 99.999999999% (11 9's)
- Availability: 99.9%
- Cost: 40-50% lower than S3 Standard (but retrieval costs apply)

#### **Use Cases:**

- ✓ Backups that need instant availability.
- ✓ Disaster recovery files that should be available but rarely accessed.

- ✓ Archived customer invoices, contracts, and compliance reports.
- ✓ Large media assets (e.g., movies, raw footage) that are accessed occasionally.

**Example:** A healthcare company storing **patient records** that need to be retained for compliance but are accessed rarely.

# **X4**S3 One Zone-IA – Lower Cost for Non-Critical Data

- Best For: Infrequently accessed data that doesn't need multi-AZ redundancy
- Durability: 99.99% (lower than Standard-IA)
- Availability: 99.5%
- Cost: 20% lower than Standard-IA (but stored in a single AZ)

#### **Use Cases:**

- ✓ Data that can be **recreated if lost** (temporary backups, logs).
- ✓ Intermediate data used in analytics or processing jobs.
- ✓ Non-critical development/test data.
- **Example:** A software company storing **daily error logs** for debugging, which are needed for a short time.

# **1** ■ **5** S3 Glacier – Low-Cost Archival Storage

- Best For: Long-term archives with infrequent access
- Durability: 99.99999999% (11 9's)
- Retrieval Time: Minutes to hours (depending on retrieval option)
- Cost: Up to 80% cheaper than Standard-IA

- **V** Use Cases:
- ✓ Long-term backups (e.g., compliance data, tax records).
- ✓ Archived videos, raw images, or sensor data for research.
- ✓ Legal documents that must be stored for years.
- ✓ Historical financial records (bank statements, audits).
- **Example:** A **government agency** storing **50 years of land records** that are rarely accessed but must be retained.

# **[6]** S3 Glacier Deep Archive – Lowest Cost for Long-Term

# **Storage**

- Best For: Data that must be stored for years but is rarely accessed
- Durability: 99.99999999% (11 9's)
- Retrieval Time: 12–48 hours
- Cost: Cheapest storage option in S3
- **Use Cases:**
- ✓ Regulatory & compliance archives (HIPAA, GDPR, financial records).
- ✓ Historical archives (scientific research, university data, old records).
- ✓ Massive-scale backups (e.g., newspaper articles, space research data).
- ✓ Cold storage of terabytes/petabytes of data.
- **Example:** A space research organization storing satellite images from past missions, which need to be retained but rarely accessed.

# **AWS S3 Policies**

AWS S3 policies define who can access your bucket and what actions they can perform. Proper policy management ensures security, compliance, and controlled access to S3 resources.

# **X** Types of S3 Policies

S3 supports multiple types of policies for controlling access:

# Bucket Policies

- Apply to an entire S3 bucket and all objects inside it.
- Used for **granting or restricting access** to users, roles, or accounts.
- Written in **JSON format** and attached directly to the S3 bucket.

# **Example Use Cases:**

- Allow only specific AWS accounts to access the bucket.
- Deny access to everyone except certain IAM roles.
- Make the **bucket read-only for anonymous users** (public read access).

# 🤧 IAM Policies

- Attach to IAM users, groups, or roles.
- Define **permissions** for interacting with **one or more S3 buckets**.

# **Example Use Cases:**

- Give developers read-write access to a specific bucket.
- Restrict access to only certain prefixes within a bucket.
- Allow **EC2 instances to upload files** to S3 using IAM roles.

# Access Control Lists (ACLs) (Legacy Method – Use Policies Instead)

- Grant permissions at the object level (not recommended for large-scale access control).
- Limited to read and write permissions.

# **Example Use Cases:**

- Share a specific **object** with a **third party**.
- Grant read access to another AWS account without modifying the bucket policy.

# S3 Access Points Policies

- Used for managing access to specific applications using access points.
- Provides fine-grained control by creating different policies per access point.

# **Example Use Cases:**

A data analytics team accesses S3 through one access point, while a
machine learning model accesses it through another, with different rules.

# Block Public Access Settings

- Prevents accidental public access to S3 buckets and objects.
- Overrides all policies that would make a bucket public.

# **Example Use Cases:**

- Ensure that **sensitive corporate data** is never exposed.
- Enforce security policies at the **account level**.

# S3 Bucket Policy Structure

An S3 bucket policy is written in JSON format and consists of:

- Statement(s): Defines what actions are allowed or denied.
- Effect: "Allow" or "Deny".
- Principal: Specifies who gets the permission.
- Action: Specifies what actions (like s3:GetObject, s3:PutObject) are allowed or denied.
- **Resource**: The bucket or object to which the policy applies.
- Condition (Optional): Adds extra restrictions like IP-based access, MFA authentication, etc.

# \*\* Popular Use Cases of AWS S3

AWS S3 is **one of the most widely used** storage services, powering applications across industries. From **big data analytics** to **media streaming**, its **scalability**, **durability**, **and security** make it an ideal choice for various use cases.

# Cloud Storage & Backup

- Businesses use S3 to store and back up critical data securely.
- Automated **versioning** helps prevent accidental deletions.
- S3 Lifecycle Policies allow data to be **archived or deleted** automatically.

**Example:** A **finance company** backs up its **transaction logs** daily, with automatic deletion of old logs after **one year**.

# Media Hosting & Streaming

- S3 serves as a content distribution storage for images, videos, and audio.
- Integrated with CloudFront for low-latency streaming.
- Works with S3 Intelligent-Tiering to optimize storage costs.
- **Example:** A **video streaming platform** like Netflix stores videos in **S3 Standard** for frequent access and **S3 Glacier** for archived content.

# Static Website Hosting

- S3 can host static websites (HTML, CSS, JavaScript).
- Supports custom domains with Route 53.
- Integrated with **CloudFront** for faster global delivery.
- **Example:** A startup **deploys its product landing page** on S3 instead of setting up a traditional web server.

# Big Data & Analytics

- Data lakes in S3 allow processing massive datasets with AWS services.
- Supports **S3 Select** for querying structured data.
- Integrated with Athena, Redshift, and EMR for analytics.
- **Example:** A **marketing company** analyzes **customer behavior data** stored in S3, using AWS Glue and Athena.

# Machine Learning & Al Data Storage

- ML models require huge datasets stored securely.
- S3 integrates with **SageMaker, TensorFlow, and PyTorch**.
- Used for image recognition, NLP, and recommendation systems.
- **Example:** An **Al-powered chatbot** stores training datasets in S3 and pulls them into **SageMaker** for model training.

# Disaster Recovery & Business Continuity

- S3 provides 99.999999999 durability, making it ideal for disaster recovery.
- Cross-Region Replication (CRR) ensures copies are available in different AWS regions.
- Works with AWS Backup & Glacier for long-term storage.

**Example:** A **banking system** keeps a **real-time replica** of customer transactions in **another region** to prevent data loss.

# Log Storage & Monitoring

- AWS services like CloudTrail and VPC Flow Logs store logs in S3.
- Security tools like **GuardDuty** analyze logs for threats.
- Logs can be queried using **Athena**.

**Example:** A **cybersecurity team** stores **application access logs** in S3 and runs **real-time threat detection** with AWS Security Hub.

# 💰 E-Commerce & Retail

- Online stores store product images, descriptions, and metadata in S3.
- S3 ensures high availability for global users.
- Works with AWS Lambda for dynamic image resizing.

**Example:** An **e-commerce website** like Amazon stores millions of **product images** and uses S3 to **serve optimized thumbnails** dynamically.

# **\*\*** Financial & Healthcare Data Storage

- S3 meets compliance standards (HIPAA, PCI-DSS, GDPR).
- Supports server-side and client-side encryption.
- Works with **Macie** to detect sensitive data automatically.

**Example:** A **hospital stores patient records** in an encrypted S3 bucket with strict IAM policies.

# → Mobile & Web App Data Storage

- S3 acts as a **centralized storage** for mobile and web apps.
- Integrates with AWS SDKs for real-time access.
- Works with **Cognito** for secure authentication.

**Example:** A social media app stores user profile pictures and videos in S3 with automatic resizing using AWS Lambda.



# Welcome to the Step-by-Step Guide on

# Deploying a Static Website with S3,

# CloudFront & Route 53!

# About This Guide

In today's digital world, fast, secure, and scalable websites are essential. AWS offers a powerful and cost-effective way to host static websites using Amazon S3, CloudFront, and Route 53. This guide will walk you through a real-world implementation of setting up a highly available, secure, and globally distributed static website step by step.

#### What You'll Learn

By following this guide, you will:

- Set up an S3 bucket for website hosting
- Configure CloudFront for global content delivery & performance optimization
- Integrate Route 53 for domain name mapping
- Enforce HTTPS security for a secure browsing experience
- Deploy a fully functional static website with AWS best practices

# Why This Setup?

- Scalability Handles thousands of requests seamlessly.
- Speed & Performance CloudFront's CDN accelerates content delivery worldwide.

- Security HTTPS encryption and access control enhance website protection.
- Cost-Effective Pay only for what you use, with minimal operational overhead.

# Let's Get Started!

This hands-on task ensures you not only **learn the theory** but also **implement a production-grade static website** using AWS services. Follow the steps carefully, and by the end, you'll have a **fully operational website with a custom domain**, **secured with SSL**, and **delivered at blazing speed across the globe!** 

Let's dive in!

# Real Time Task on Step By Step Implementation of S3 Static Website with Cloudfront and Route53

# **Step 1: Create and Configure the S3 Bucket**

#### 1.1 Create the S3 Bucket

- 1. Log in to the AWS Management Console and navigate to the S3 service.
- 2. Click Create bucket.
- In the **Bucket name** field, enter a unique name (e.g., my-example-bucket).
- Under Access Control List (ACL) settings, select the option that enables
   ACLs (if not already enabled).

#### 5. Block Public Access Settings:

- o Uncheck Block all public access.
- Confirm your intent if prompted.

#### 6. **Bucket Versioning:**

Enable Bucket versioning to keep versions of objects.

#### 7. Click Create bucket.

#### 1.2 Upload Website Content to the S3 Bucket

- 1. Open the newly created bucket.
- 2. Click Upload.
- Upload your website file(s) at a minimum, an index.html file containing your website content.
- 4. After uploading, select the uploaded object(s) and configure permissions:
  - o Grant **public-read** access so that the content is available publicly.
- 5. Confirm that your object now shows public-read permissions.

#### 1.3 Enable Static Website Hosting on the S3 Bucket

- 1. In your S3 bucket, click on the **Permissions** tab.
- 2. Scroll down to the **Static website hosting** section.
- 3. Select **Enable**.
- 4. For the **Index document**, enter index.html.
- 5. (Optionally, specify an error document.)
- 6. Click Save changes.

# **Step 2: Create a CloudFront Distribution**

#### 2.1 Launch a New CloudFront Distribution

- 1. Open the AWS Management Console and navigate to CloudFront.
- 2. Click Create Distribution.
- 3. Under **Web** (the distribution method for web content), click **Get Started**.

#### 2.2 Configure the Origin Settings

#### 1. Origin Domain Name:

 From the dropdown, select your S3 bucket (it will appear as something like my-example-bucket.s3.amazonaws.com).

# 2. Origin Access:

- Under Origin Access, choose Legacy OAI and select or create an
   Origin Access Identity (OAI) that grants CloudFront permission to access your bucket.
- (Ensure that your S3 bucket policy allows access from this OAI if needed.)

#### 2.3 Configure Viewer Protocol Policy

1. In the **Default Cache Behavior Settings**, find **Viewer Protocol Policy**.

Select Redirect HTTP to HTTPS to ensure all requests use a secure connection.

#### 2.4 Enable Web Application Firewall (WAF) (Optional)

 Under Web Application Firewall (WAF), choose Enable security protections if you want to attach a WAF for additional security.

#### 2.5 Configure Alternate Domain Names and SSL Certificate

- In the Settings section, locate Alternate Domain Names (CNAMEs).
- 2. Enter your domain name with the www prefix (e.g., www.example.com).
- 3. Under SSL Certificate, select Custom SSL Certificate (example.com).
  - Choose the ACM certificate that covers www.example.com.

#### 2.6 Set Default Root Object

1. For **Default Root Object**, enter index.html.

#### 2.7 Create the Distribution

- 1. Review all your settings.
- 2. Click Create Distribution.
- 3. Wait for the distribution to deploy (this might take up to 20–30 minutes).

4. Once deployed, copy the **Distribution Domain Name** (e.g.,

d123456abcdef8.cloudfront.net); you will use this in Route 53.

# **Step 3: Configure Route 53 DNS**

#### 3.1 Create a Record to Point to CloudFront

- 1. Open the AWS Management Console and navigate to Route 53.
- 2. Go to **Hosted Zones** and select your hosted zone (e.g., example.com).
- 3. Click Create Record.
- 4. In the **Record name** field, enter www (this creates www.example.com).
- 5. Record Type: Choose A IPv4 address.
- 6. Alias: Enable the alias option.
- 7. Alias Target:
  - In the dropdown, select the CloudFront distribution you created (e.g., d123456abcdef8.cloudfront.net).
- 8. Click Create Record.
- 9. Wait until the record status shows as **insync**.

# **Step 4: Verify the Setup**

#### 1. **DNS Propagation:**

Open a terminal or use an online DNS checker and run:

nslookup www.example.com

Ensure it resolves to your CloudFront distribution.

#### 2. Access the Website:

Open your web browser and navigate to

https://www.example.com.

 You should see the content of your index.html file (hosted in your S3 bucket).

# **Summary: What This Scenario Does**

#### • S3 Bucket:

- Creates a bucket that hosts your static website content (e.g., index.html).
- Configured for static website hosting with public-read permissions.

#### • CloudFront Distribution:

 Serves as a global content delivery network (CDN) that caches and delivers your website content.

- Redirects HTTP requests to HTTPS and uses your ACM certificate for secure connections.
- Uses an Origin Access Identity (OAI) to securely access the S3 bucket.

#### Route 53 DNS:

- Creates a DNS record (www.example.com) that points to the CloudFront distribution.
- Ensures that when users visit your domain, they are routed to your
   CloudFront distribution, which in turn serves the content from S3.

#### Overall Outcome:

 A secure, highly available, and globally distributed static website hosted on S3, served through CloudFront, and accessed via your custom domain (www.example.com).

# What Does This Scenario Do? (In Simple Terms)

This scenario sets up a static website using AWS services **S3**, CloudFront, and Route **53**, making it secure, fast, and accessible via a custom domain.

Here's how it works step by step:

#### 1. Store Website Files in S3 (Simple Storage Service)

- You create an **S3 bucket** and upload your website files, including an index.html file.
- You enable **public access** so the website can be viewed by anyone.
- You enable static website hosting in S3.

#### 2. Use CloudFront (Content Delivery Network) for Faster & Secure Delivery

 You create a CloudFront distribution that fetches your website content from S3.

- CloudFront caches the content across multiple locations worldwide, making it load faster.
- It enforces **HTTPS**, ensuring your website is secure.
- You link an SSL certificate (ACM) for a custom domain (www.example.com).

#### 3. Route the Custom Domain to CloudFront Using Route 53

- You create a DNS record in Route 53 that maps www.example.com to the CloudFront distribution.
- This allows users to visit your site using your own domain name instead of an AWS URL.

#### **Final Outcome:**

When someone types www.example.com in their browser:

- Route 53 directs the request to CloudFront.
- CloudFront fetches and serves the website files from S3.
- The website loads quickly, securely (HTTPS), and globally with reduced latency.

This setup makes your website:

- Fast (via CloudFront caching)
- Secure (via HTTPS & SSL certificate)
- Scalable (AWS handles traffic spikes)
- Cost-effective (S3 + CloudFront is cheaper than traditional hosting)

# Welcome to the Real-Time Task on S3

# Policies, Access Points, and S3FS!

In this hands-on task, we will explore the powerful capabilities of Amazon S3 by implementing key security and access management features. This step-by-step guide will help you gain practical experience with:

- S3 Bucket Policies Secure and manage permissions at the bucket level.
- ✓ Pre-Signed URLs Enable temporary and secure access to S3 objects.
- S3 Access Points Implement fine-grained access control for different users.
- S3FS Integration Mount an S3 bucket as a filesystem on an EC2 instance.

#### What You'll Achieve:

- Launch and configure an Ubuntu EC2 instance.
- Create and configure an S3 bucket with security best practices.
- Generate and test pre-signed URLs for secure temporary access.
- Implement S3 Access Points to manage access for different IAM users.
- Mount your S3 bucket as a filesystem on an EC2 instance using s3fs.

This task will provide you with **real-world AWS experience** and enhance your understanding of secure storage management in the cloud. By the end, you will have a fully functional S3 setup integrated with IAM users, access points, and EC2.

Let's get started and dive into AWS storage security and access management! 🚀

# Real Time Task on Step by Step Implementation of S3 Policies , Access Points and S3FS

#### Part 1: Launch an Ubuntu EC2 Instance

#### 1. Launch an EC2 Instance:

- Open the AWS Management Console and navigate to EC2.
- Click Launch Instance.
- Choose an **Ubuntu AMI** (for example, Ubuntu 20.04 LTS).
- Select the instance type (e.g., **t2.micro** for testing).
- Under Configure Instance Details, select your desired VPC (ensure its security group allows all traffic from anywhere) and choose a Public Subnet.

- o Enable Auto-assign Public IP.
- Under Add Tags, add a tag such as:

■ Key: Name

■ Value: Testing

- Proceed to Configure Security Group ensuring that necessary ports
   (e.g., SSH on port 22) are open.
- Review and launch the instance.
- Select your **PEM key pair** when prompted.
- Wait for the instance to be in the **running** state.

# Part 2: Create and Configure an S3 Bucket

#### 2.1 Create an S3 Bucket and Upload Files

- 1. Create the Bucket:
  - o In the AWS Console, navigate to **S3**.
  - Click Create bucket.
  - Bucket Name: Enter a unique name (for example, my-example-bucket).
  - Leave the default settings and click Create bucket.

#### 2. Upload Files:

- Open the newly created bucket.
- Click **Upload** and add some files (for example, images) including an index.html file.
- Complete the upload process.

#### 2.2 Modify Bucket Permissions to Allow Public Read (Temporarily)

#### 1. Disable Block Public Access:

- o In your bucket, go to the **Permissions** tab.
- Click on **Edit** in the **Block public access (bucket settings)** section.
- Uncheck the option Block all public access.
- Click Save changes (confirm if prompted).

#### 2. Add a Bucket Policy for Public Read:

o In the same **Permissions** tab, scroll to **Bucket Policy**.

Click **Edit** and paste the following JSON (replace

```
arn:aws:s3:::my-example-bucket/* with your bucket ARN pattern):
{
    "Version": "2012-10-17",

    "Statement": [
```

```
"Sid": "PublicReadGetObject",

"Effect": "Allow",

"Principal": "*",

"Action": "s3:GetObject",

"Resource": "arn:aws:s3:::my-example-bucket/*"
}
]
```

• Click **Save changes**.

#### 3. Test Public Access:

- Copy the URL of one of your files (e.g.,
   https://my-example-bucket.s3.amazonaws.com/index.h
  - tml) and paste it in a browser.
- o Confirm that you can view the file.

#### 4. Revert Public Access Settings:

- $\circ\ \ \,$  For security reasons, after testing, go back to the  $\mbox{\bf Permissions}$  tab.
- o Re-enable **Block all public access**.
- o Delete the bucket policy you just created.

# Part 3: Generate a Pre-signed URL

#### 1. Generate a Pre-signed URL:

- In the S3 Console, navigate to your bucket.
- o Select any file (e.g., one of your images).
- Click on Actions and choose Share with pre-signed URL.
- Specify the expiration time (e.g., 60 minutes).
- Click Create pre-signed URL.

#### 2. Test the Pre-signed URL:

 Copy the pre-signed URL and paste it into a browser to ensure you can access the file.

#### 3. Clean Up:

o Optionally, delete all files in this bucket if they were only for testing.

# **Part 4: Configure S3 Access Points**

#### 4.1 Create Folders for Granular Access

- 1. Open your S3 bucket.
- 2. Click Create folder and create two folders:

- o folder1
- folder2
- (The scenario is designed such that you want to give different permissions to different developers for each folder.)

#### 4.2 Create an IAM User for Developer Access

- 1. In the AWS Console, navigate to IAM.
- 2. Go to Users and click Add users.
- 3. Username: Enter developer1.
- Access Type: Choose AWS Management Console access (or only programmatic access if you plan to use CLI).
- 5. Set a **custom password** and create the user without attaching any policies.
- 6. Click Create user.

#### 4.3 Create an S3 Access Point

- 1. Go to your S3 bucket.
- 2. Click on the Access Points tab.
- 3. Click **Create access point**.
- 4. Name: Enter accesspointdev1.
- 5. **Network Origin:** Select **Internet**.
- 6. Leave other settings as default.

7. Click Create access point.

{

#### 4.4 Update Bucket Policy for Access Point Usage

1. Go back to your S3 bucket's **Permissions** tab.

Under **Bucket Policy**, add (or update) a policy that allows AWS users to perform actions on this bucket only when using a data access point from your AWS account. For example:

```
"Version": "2012-10-17",
"Statement": [
  {
    "Sid": "AllowActionsViaAccessPoint",
    "Effect": "Allow",
    "Principal": "*",
    "Action": "s3:*",
    "Resource": "arn:aws:s3:::my-example-bucket/*",
    "Condition": {
      "StringEquals": {
```

2. Click **Save changes**.

# 4.5 Attach an Access Point Policy for Developer1

- 1. In the S3 bucket, click on the **Access Points** tab.
- 2. Select the access point you created (e.g., accesspointdev1).
- 3. Click on the **Permissions** section for the access point.

Edit the **Access Point Policy** to allow the IAM user developer1 to perform any action on all objects under **folder1**. For example:

```
{
   "Version": "2012-10-17",
   "Statement": [
```

```
{
      "Sid": "Developer1FullAccessToFolder1",
      "Effect": "Allow",
      "Principal": {
        "AWS":
"arn:aws:iam::<YOUR_AWS_ACCOUNT_ID>:user/developer1"
      },
      "Action": "s3:*",
      "Resource": [
        "arn:aws:s3:::my-example-bucket/folder1/*"
    }
}
  4. Click Save.
```

#### 4.6 Test Developer1 Access via S3 Access Point

On your local machine (or via Git Bash), configure the AWS CLI as **developer1**: aws configure

1.

 Enter the credentials (Access Key ID, Secret Access Key) for developer1.

Create an empty file on your local machine:

```
echo "Test file for access point" > testfile.txt
```

Copy the file to your S3 bucket using the access point URL. The S3 URI will follow a format similar to:

```
aws s3 cp testfile.txt
s3://<accesspointdev1-identifier>/folder1/testfile.txt
```

2.

- The exact S3 access point URL format can be found in the Access
   Point details.
- 3. Verify that the file is uploaded successfully into the folder1 directory.

#### Part 5: Use s3fs to Mount the S3 Bucket

#### 5.1 Create an IAM User with Full S3 Access

- 1. In the AWS Console, navigate to IAM.
- 2. Create a new IAM user named s3full with Programmatic Access.
- 3. Attach the **AmazonS3FullAccess** policy to this user.
- 4. After creation, note down the **Access Key ID** and **Secret Access Key**.

#### **5.2** Configure the EC2 Instance for s3fs

Log in to your Ubuntu EC2 instance (launched earlier) via SSH:

```
ssh -i your-key.pem ubuntu@<Ubuntu_Public_IP>
```

Update the package index:

sudo apt update

Install required packages:

sudo apt install s3fs unzip awscli -y

Configure AWS CLI with the credentials of the s3full user:

aws configure

- Enter the Access Key ID and Secret Access Key for s3full.
- o Set the default region as appropriate.
- Set output format to json (or your choice).

Create a directory to mount the S3 bucket:

Mount the S3 bucket using s3fs:

```
s3fs my-example-bucket ~/mys3 -o use_cache=/tmp
```

• Replace my-example-bucket with your S3 bucket name.

Verify the mount:

You should see the S3 bucket mounted at ~/mys3.

# **Summary: What This Scenario Does**

#### S3 Bucket Setup:

- A bucket is created and files (such as images, an index.html) are uploaded.
- Public access is enabled temporarily by modifying the bucket policy, and later reverted.
- A pre-signed URL is generated to provide temporary access to an object.

#### S3 Access Points:

- Two folders (folder1 and folder2) are created in the bucket.
- An S3 Access Point (accesspointdev1) is created to allow granular access.
- A bucket policy is configured so that actions on the bucket are allowed only via the Access Point, and a separate access point policy is set to give a specific IAM user (developer1) permissions on folder1.

#### • Testing Access via CLI:

The developer1 IAM user uses AWS CLI to upload a file to folder1
 using the Access Point URL, verifying that the policy is working.

#### Using s3fs:

- o An IAM user (s3full) with full S3 access is created.
- An EC2 Ubuntu instance is configured with s3fs and awscli, and the
   S3 bucket is mounted as a local filesystem for further operations.

This scenario demonstrates different ways to manage S3 access:

- Bucket policies for public object access.
- **Pre-signed URLs** for temporary secure sharing.

- **S3 Access Points** for granular, per-folder permissions.
- s3fs for mounting S3 buckets on Linux systems for file system-like access.

# Understanding the S3 Policies, Access Points, and S3fs Scenario Task

This scenario is a **step-by-step implementation** of different ways to **access**, **secure**, **and mount an Amazon S3 bucket** while implementing IAM permissions and S3 Access Points. Let's break down what this task actually achieves.

#### 1. Launch EC2 Instance with Ubuntu AMI

We start by launching an EC2 instance that will be used to interact with S3. This is necessary because we need a server to perform AWS CLI operations and later mount S3 as a filesystem.

# 2. Basic S3 Bucket Access and Public Permissions Testing

Here, we create an S3 bucket and explore different ways of accessing objects stored in it.

#### **Steps and Purpose:**

- Create an S3 bucket This is a storage container for files (images, logs, backups, etc.).
- 2. **Upload files (e.g., images) to the S3 bucket** This helps us test object accessibility.
- 3. Modify bucket permissions:

- Disable "Block all public access" This allows public access to objects.
- Add a bucket policy We create a policy that allows anyone (public) to retrieve objects (s3:GetObject action).
- Test access via URL Now, we try to access an object's URL from a browser to verify it's publicly accessible.
- Re-enable "Block all public access" We remove public access to secure the bucket.

#### What this achieves:

- We learn how to make an S3 bucket public (for testing) and later revert it to private.
- We see how **Bucket Policies** work for granting public access.

# 3. Pre-Signed URL for Temporary Access

Instead of making files public, we generate a pre-signed URL for an object.

#### **Steps and Purpose:**

- 1. Select an S3 object → Generate a Pre-Signed URL
  - This URL is time-limited (e.g., expires after 10 minutes or 1 hour).
  - The object remains private, but anyone with the pre-signed URL can access it temporarily.
- 2. **Test access in a browser** Open the pre-signed URL and verify access.
- 3. **Delete all files from the bucket** Cleanup before moving to the next step.

#### What this achieves:

- Pre-signed URLs allow temporary access to private S3 objects without changing bucket policies.
- A common use case is **sharing private files securely for a limited time**.

# 4. Implementing S3 Access Points for Developer-Specific Access

S3 Access Points help manage permissions for **specific users** and **folders** within a bucket.

#### Scenario:

We assume **two developers**, and we want to **grant each developer access to only their respective folder**.

#### **Steps and Purpose:**

#### (a) Create Folders in the S3 Bucket

 Folder1 and Folder2 – Representing different project areas for different developers.

#### (b) Create IAM User (Developer1)

No permissions are granted initially.

#### (c) Create an S3 Access Point

- We create an Access Point (accesspointdev1) for the S3 bucket.
- This allows controlled access only through the Access Point.

#### (d) Apply an Access Point-Specific Policy

- Modify S3 Bucket Policy:
  - Ensures that all access to the bucket happens only via the Access
     Point.
  - This prevents direct S3 access via bucket URLs.
- Modify Access Point Policy:
  - Grants Developer1 access to Folder1 only.

#### (e) Developer1 Uploads a File to S3 via Access Point

- Developer1 configures AWS CLI (aws configure).
- Uploads a file to Folder1 using the Access Point URL.

#### What this achieves:

- Restricts access to specific users based on Access Points.
- Ensures access is managed centrally via the Access Point, not the bucket.
- Prevents unintended access to other folders.

# 5. Mounting S3 Bucket as a Filesystem Using S3fs

Finally, we mount the S3 bucket as a local filesystem on an EC2 instance.

#### **Steps and Purpose:**

#### (a) Create IAM User (s3full) with AmazonS3FullAccess

This user has full access to S3.

#### (b) Install Required Packages on the EC2 Instance

- **s3fs** A tool that allows mounting S3 as a filesystem.
- AWS CLI To interact with S3.
- **Unzip** Helps extract files if needed.

#### (c) Configure AWS Credentials on EC2 (aws configure)

• We enter the Access Key & Secret Key for s3full.

#### (d) Mount the S3 Bucket to a Local Directory

- Create a directory (mys3/) to serve as the mount point.
- Run s3fs to mount the bucket.
- **Verify with df** -h The bucket should appear as a mounted drive.

#### What this achieves:

- We can now treat the S3 bucket like a local filesystem.
- Developers can **read**, **write**, **and manage files** on S3 using standard Linux commands (1s, cp, mv, etc.).
- Useful for backup storage, logging, and cloud file management.

# Wrapping Up: AWS S3 Service

# Comprehensive Guide 6

We started this journey by understanding the **theoretical concepts** behind S3 Policies, Access Points, and S3FS, ensuring you have a solid foundation before diving into the hands-on part.

Then, we **stepped into real-world implementation** with a practical approach:

- Launching an Ubuntu EC2 Instance to serve as our workspace
- Creating and Configuring an S3 Bucket with proper permissions
- Managing S3 Access Control with Policies and Access Points
- Generating Pre-Signed URLs for temporary access
- Mounting S3 Storage with s3fs to use it like a local drive

This structured approach helped you learn, apply, and validate your skills with AWS S3 in a real-world scenario. By completing this task, you've gained hands-on experience in securely managing cloud storage using AWS best practices.

#### What's Next? 🚀

The learning never stops! Follow me for more real-time DevOps and Cloud tasks daily. Every day, we explore new, industry-relevant projects that will take your skills to the next level.

Stay consistent, keep practicing, and get ready for the next challenge!

Until next time—Happy Learning & Keep Building! 🔥 🚀

